

Constant-Time Discrete Gaussian Sampling

Angshuman Karmakar¹, Sujoy Sinha Roy¹, Oscar Reparaz²,
Frederik Vercauteren³, and Ingrid Verbauwhede³

Abstract—Sampling from a discrete Gaussian distribution is an indispensable part of lattice-based cryptography. Several recent works have shown that the timing leakage from a non-constant-time implementation of the discrete Gaussian sampling algorithm could be exploited to recover the secret. In this paper, we propose a constant-time implementation of the Knuth-Yao random walk algorithm for performing constant-time discrete Gaussian sampling. Since the random walk is dictated by a set of input random bits, we can express the generated sample as a function of the input random bits. Hence, our constant-time implementation expresses the unique mapping of the input random-bits to the output sample-bits as a Boolean expression of the random-bits. We use bit-slicing to generate multiple samples in batches and thus increase the throughput of our constant-time sampling manifold. Our experiments on an Intel i7-Broadwell processor show that our method can be as much as 2.4 times faster than the constant-time implementation of cumulative distribution table based sampling and consumes exponentially less memory than the Knuth-Yao algorithm with shuffling for a similar level of security.

Index Terms—Knuth-Yao, constant-time sampling, lattice-based cryptography, discrete Gaussian sampling

1 INTRODUCTION

PUBLIC-KEY cryptography (PKC) eliminated one serious drawback of otherwise highly efficient symmetric-key cryptography, namely requirement of key establishment among all the communicating parties or the existence of a central key distribution authority. PKC overcomes this problem by disseminating a global public-key and a secret private-key. The security of such public-key cryptosystems are assured by underlying computationally hard problems. Since the discovery of the Diffie-Hellman [1] key exchange protocol, the popularity and utility of PKC has grown steadily over the past few decades. Currently, primitives derived from RSA and ECC are used extensively for public-key cryptography on a wide range of devices. In comparison to symmetric-key cryptography, the major drawbacks of PKC are larger key sizes and slower running time. To get the best of both worlds, contemporary security protocols use both schemes in tandem for highly efficient and secure digital security solutions.

Unfortunately, large-scale quantum computers running Shor's [2] and Proos-Zalka's [3] algorithms can solve the underlying hard problems of RSA and ECC. In this scenario, lattice-based PKC [4],[5] has become an attractive choice to

provide digital security in the post-quantum world. The confidence in security of such schemes arises from the fact that unlike RSA and ECC, there is no known algorithm that can use quantum computers to efficiently solve the underlying hard problems of lattice-based cryptography. Hard lattice problems like Learning with errors (LWE) [4], Short Integer Solutions [5] and their ring equivalents R-LWE, R-SIS [6], [7] are some of the prominent choices to build various lattice-based cryptography protocols. In fact, there exists a wide variety of cryptography primitives that can be built on top of these problems. Examples are, digital signature schemes [8], [9], [10], public-key encryption [6], [11], key-exchange protocols [12], [13], identity-based encryption [14], [15], [16], [17], fully homomorphic encryption [18], [19], [20], [21] etc. The other features which make lattice-based cryptography a suitable alternative are, proven worst case to average case reduction of lattice problems and somewhat simpler operations than other PKC schemes.

LWE is a system of *approximate linear equations* with the secret key being the solution of the system. LWE uses noise to hide its secret values, without which the system can be easily solved using Gaussian elimination. This noise is typically sampled from a discrete Gaussian distribution. Sampling from such a distribution generally involves either storing a large table of precomputed values or computing the exponential function to a very high precision (binomial sampling in BLISS [8] avoids calculating exponential function by storing a precomputed table and rejection sampling). Hence, Gaussian sampling accounts for a non-negligible share of resources and computation time in a lattice-based cryptography implementation. For example, in the case of BLISS [8] and Lyubashevsky's [22], [23] signature scheme, the Gaussian sampling alone takes about 35 and 50 percent of the total running time of the signature algorithms respectively. Since the dawn of lattice-based cryptography, a lot of research has been performed to reduce the storage and computational

- A. Karmakar, S. Sinha Roy, O. Reparaz and I. Verbauwhede are with the KU Leuven ESAT/COSIC and imec, Kasteelpark Arenberg 10, Leuven-Heverlee B-3001, Belgium. E-mail: {angshuman.karmakar, sujoy.sinharoy, oscar.reparaz, frederik.vercauteren, ingrid.verbauwhede}@esat.kuleuven.be.
- F. Vercauteren is with the KU Leuven ESAT/COSIC and IMEC, Kasteelpark Arenberg 10, Leuven-Heverlee B-3001, Belgium and also with the Open Security Research, FangDa Building 704, Kejinan-12th, Nanshan, Shenzhen 518000, China.

Manuscript received 31 May 2017; revised 8 Feb. 2018; accepted 14 Feb. 2018.
Date of publication 11 Mar. 2018; date of current version 16 Oct. 2018.

(Corresponding author: Angshuman Karmakar.)

Recommended for acceptance by Ç. K. Koç, Z. Liu, and P. Longa.

For information on obtaining reprints of this article, please send e-mail to: reprints@ieee.org, and reference the Digital Object Identifier below.

Digital Object Identifier no. 10.1109/TC.2018.2814587

overhead of sampling [8], [24], [25], [26], [27], [28]. Yet the discrete Gaussian sampler is arguably most vulnerable to side channel attacks in a lattice-based cryptography implementations. Currently, as lattice-based cryptography is becoming more efficient and being implemented in a wide variety of devices, it is imperative to make the sampling secure against side channel attacks. Different methods have been proposed to make the sampling efficient and resource friendly, but there is a lack of research to make the Gaussian sampling secure against side-channel attacks.

This was not a cause for a serious concern as there was no attack available that could efficiently exploit the side channel leakage information against the cryptosystem. Recently, Bruinderink et al. [29] has described a very effective side channel attack on the software implementation of the BLISS digital signature scheme. They exploited the irregular cache memory access pattern of the Gaussian sampler to extract information about the secret key and used the LLL lattice reduction [30] algorithm on this information to mount the attack. The attack requires only 450 signatures to recover the secret key of the BLISS signing algorithm. A prior work by Roy et al. [31] analyzed timing and power leakage from discrete Gaussian sampling in the context of public-key encryption and proposed a countermeasure based on random shuffling. Pessl [32] analyzed the shuffling-based countermeasure in detail and proposed a profiled side channel attack that can recover the key by observing only 7,000 signatures. The author proposed to use Gaussian convolution in conjunction with shuffling to increase side channel resistance. Constant-time table scanning [33] or the random shuffling methods can eliminate or mitigate the side channel leakage, but they come with a performance cost. We also note that, due to the side channel vulnerability of discrete Gaussian sampling, currently there is a trend to design lattice-based cryptography schemes that do not use Gaussian sampling in the performance critical parts of the schemes [9], [10]. These schemes however require more arithmetic operations and a larger modulus for security.

1.1 Our Contributions

In this paper, we describe a method to sample from a discrete Gaussian distribution securely. Our contributions can be summarized as follows.

- Almost all of the currently known ‘efficient’ sampling algorithms use data-dependent branches and hence their execution time varies depending on the data. In this work, we avoid any data-dependent branching to achieve constant time execution. More precisely, we analyze the Knuth-Yao discrete Gaussian sampling [27] and observe a unique mapping between the output sample values and the input random bits of the sampling algorithm. We utilize this observation to express the output sample values as a Boolean function of the input random bits. During sampling, each of these Boolean functions are evaluated in constant-time to generate each sample, thus making the sampling procedure a constant-time operation. This is described in Section 3.1.
- In Section 3.2, we show how we can exploit a bit-slicing methodology to generate samples in batches.

This increases the throughput by the order of the word length of a processor. This enhancement in performance is achieved by carefully tweaking the way random input bits are stored and utilizing bit-wise operators and the wide data path of modern processors.

- In Section 3.3, we analyze the constant-time behaviour of our sampler in practice. We analyze the effect of a smart compiler on the constant-time execution of our sampling routine. We also provide constant-time behaviour of our method on different levels of design abstraction at algorithm level, instruction level, register-transfer (RT) level, logic level etc.
- In Section 4, we compare our method to other secure discrete Gaussian samplers for a similar level of security. We provide an experimental comparison of run times using a C implementation on a Intel i7-6600 Broadwell processor. Additionally, we describe a method to split a discrete Gaussian distribution with large standard deviation into many smaller discrete Gaussian distributions with smaller standard deviation.
- In Sections 4.2.2 and 4.2.3, we also provide results of our sampling method implemented in hardware using FPGA and in software using AVX vector instructions utilizing the wider data path.
- Finally in Section 5, we provide a side channel analysis of our sampling algorithm.

2 DISCRETE GAUSSIAN SAMPLING

In this section, we provide a brief discussion on discrete Gaussian sampling and different methods to generate samples from such distribution.

2.1 Definition

The probability distribution function $D_{\mathbb{Z},\sigma}$ of a discrete Gaussian distribution defined over \mathbb{Z} with mean $\mu = 0$ and standard deviation σ is defined as,

$$D_{\mathbb{Z},\sigma}(X = z) = \frac{1}{S} e^{-z^2/2\sigma^2}.$$

Here X is a random variable defined over \mathbb{Z} and S is the normalization constant, defined as,

$$S = \sum_{x=-\infty}^{\infty} e^{-x^2/2\sigma^2} \approx \sigma\sqrt{2\pi}.$$

To generate samples over \mathbb{Z} , it is sufficient to generate samples over \mathbb{Z}^+ and use a single random bit to determine the sign due to the symmetry of discrete Gaussian distribution across its mean.

Ideally, the support of a Gaussian distribution has a range $(-\infty, \infty)$, but in most practical applications it is neither feasible nor required to generate samples from this range. Instead, *tail-cut* factor τ is used to generate samples from a smaller interval $[-\tau\sigma, \tau\sigma]$, ignoring other values beyond this interval that have very low probabilities of occurrence. Also, as the probabilities of $D_{\mathbb{Z},\sigma}(x)$, $x \in [-\tau\sigma, \tau\sigma]$ are real numbers, their binary expansions can be infinitely long. Therefore in practice, the probabilities are calculated only up to a certain precision λ depending upon the requirement of the application. For most

lattice-based cryptographic applications the values of τ and σ are chosen as 12 and 128 respectively, such that the generated samples are statistically very close to the ideal Gaussian distribution. Traditionally, statistical distance had been used in the literature to measure this closeness. But recently the work of Bai et al. [34] has shown that the value of τ can be reduced to as low as 6 using the Rényi divergence as the closeness measure for some applications.

2.2 Sampling from a Discrete Gaussian

Sampling from a continuous Gaussian distribution has a wide range of applications in different fields of natural science, social sciences, mathematics, and engineering. Hence, it has been studied extensively for long time. Sampling from a discrete Gaussian distribution is a comparatively less studied topic. Since the start of their use in lattice-base cryptography, several methods have been proposed to sample from a discrete Gaussian distribution. Some of them are rejection sampling [25], cumulative distribution table (CDT) based sampling [24], discrete Ziggurat sampling [26], Knuth-Yao sampling [27], and Bernoulli sampling [8]. Among these methods, the rejection sampling does not require any storage of precomputed tables but requires many random bits and many repetitions. Hence, it does not perform very well in practice. Almost all other methods use precomputed tables and binary search for efficient sampling. Here we discuss CDT sampling and Knuth-Yao sampling as these two methods can be more efficiently [35] instantiated as leakage-resistant sampling algorithm than others.

2.2.1 The CDT Based Sampling:

The CDT based sampling precomputes a cumulative distribution function (CDF) table T for $i \in [-\tau\sigma, \tau\sigma]$ according to the given discrete Gaussian distribution with λ bits of precision, such that $T[i+1] - T[i] = D_\sigma(i)$. The sampling phase of the algorithm is basically a search operation on the CDF table T . First, a random $r \in [0, 1)$ is generated then the table T is searched to find an s , such that $T[s+1] \geq r > T[s]$. If such an s is found, it is returned as the sample. To reduce the storage requirement for the sampling, only the interval $[0, \tau\sigma] \in \mathbb{Z}^+$ needs to be searched, as explained in Section 2.1. To improve the efficiency, binary search or improved versions of binary search such as binary search with guide table [36], [37] are used. However, in such methods the irregular table access pattern of binary search makes the sampling process vulnerable to cache-timing attacks which was used by Bruinderink et al. [29].

2.2.2 The Knuth-Yao Sampling

The Knuth-Yao [38] sampling algorithm was proposed to generate samples from any source with a known probability distribution. The sampling algorithm uses a rooted binary tree which in this context is also known as a discrete distribution generation (DDG) tree. The DDG tree is constructed from the probability matrix, which is a matrix constructed from the samples in the support of the distribution and their corresponding binary expanded probabilities up to a certain precision. The probability matrix and the DDG tree are related as follows: the number of leaf nodes in the DDG tree at the i th level is equal to the Hamming weight of the i th




Fig. 1. A Probability matrix and the DDG tree corresponding to it. The random bits $\{0, 1\}$ are used to traverse the tree starting from the root.

column of the probability matrix. Each leaf node of the DDG tree corresponds to a sample in the sample space. An example of the probability matrix and the corresponding DDG tree is shown in Fig. 1 for an arbitrary toy distribution with a sample space S consisting of only four samples.

The sampling operation is a random walk on the DDG tree. The random walk starts from the root and at each non-leaf node a random bit is used to determine the direction of the random walk in the left or right sub-tree. The random walk stops when it hits a leaf node and the corresponding sample is returned. Here, the non-constant running time and branching during the random walk expose the cache vulnerability of the sampling operation.

Dwarakanath and Galbraith [27] first adapted the Knuth-Yao algorithm to sample from a discrete Gaussian distributions. Their work was later extended by Roy et al. [28] with a more simplistic design methodology and reduced memory requirement. We refer the interested readers to their work for further details.

2.3 Previous Works

As noted in Section 1, there has not been much research on the construction of constant-time Gaussian samplers, largely because of non-existence of efficient side-channel attacks. However, the existing non constant-time Gaussian samplers can be used for secure Gaussian sampling by applying some simple countermeasures. In this section we briefly revisit them.

2.3.1 Constant-Time Table Access

The table-based Gaussian samplers use binary search for efficiency, which also makes them vulnerable to timing attacks. These algorithms can be converted to secure sampling algorithms by replacing the binary search with constant-time linear search of the whole table. This removes the cache-weakness of the binary search. But this countermeasure does not perform very well in practice. Linear search of the whole table incurs a significant penalty in performance. Bos et al. [33] used this method for a leakage-resistant Gaussian sampling in their key-exchange scheme.

2.3.2 Shuffling

Constant-time table access for Knuth-Yao sampling is more complicated and inefficient than the other table-based sampling methods. Roy et al. [31] proposed a method to

mitigate the problem of side-channel leakage of the Knuth-Yao sampler using extra memory. Their method caches the first k columns of the probability matrix in a table with 2^k entries. The table entries are either a sample value or an intermediate position in the DDG tree. The sampling operation of this algorithm can be divided in a secure and a non-secure part. In the secure part, the algorithm generates a k bit random index and looks for the entry in the table. If the entry is a sample value, then it is returned. In the non-secure part, if the table entry holds a position in the DDG tree, then a random walk is commenced from that position to find a sample. In this scenario, the algorithm leaks the absolute values of the samples due to the difference in timing to find a sample. As a second countermeasure, the authors suggest a random permutation of the leaked and non-leaked samples after the sampling to obfuscate the locations of the samples from the attacker. Also, as the security of this method depends on the number of columns cached, the memory requirement increases exponentially with an increase in the levels of security.

For the sake of completeness, we should also mention here that it is possible to create a constant-time Knuth-Yao sampler by scanning the probability matrix table fully for each sample using Alg.1 in [28]. But, in this case the number of items to be scanned is very large and the performance becomes poor. We discuss this further in Section. 4.

2.3.3 Fixed Step Binary Search

Howe et al. [35] proposed a fixed step binary search for secure Gaussian sampling. In their proposal, the binary search always runs for $O(\log(n))$ steps where n is the size of the table, irrespective of whether the sample has been found in a previous step or not. While this method may work for some specific platforms, it is not a generic solution for constant-time sampling. The binary search will leak secrets on a wide variety of platforms due to irregular memory access patterns.

3 CONSTANT-TIME KNUTH-YAO SAMPLING

In this section, we analyze the Knuth-Yao sampling algorithm. We describe our observation on correlation between samples and input random bits. Based on this observation we propose a constant-time Knuth-Yao discrete Gaussian sampling. We also propose an optimization scheme to increase the throughput of our sampling algorithm. We conclude the section with an analysis of our constant-time sampler on different levels of execution.

Choice of Sampling Algorithm. At this point, we describe our rationale for choosing the Knuth-Yao algorithm for constant-time sampling. During our initial investigation for a constant-time Gaussian sampler, we found two possible methods to devise a constant-time sampler. One is the simple constant-time linear table search, the other method is to express each sample as a function of input bits and then execute the function in constant-time to calculate each sample. The former method is a well known method and has been used before. But, there is no precedence in the literature of the latter method for constant-time sampling. For the second method we require a well defined mapping from the input random bits to the output samples. We found that




Fig. 2. Mapping $f^i : \{0,1\}^n \rightarrow \{0,1\}$ from set of random bit strings to the bits of samples.

due to the random-walk nature of Knuth-Yao sampling it is easier to find such a mapping (explained later) and hence a function that can be executed efficiently to find samples from the input random bits. We also stress that we do not claim that such an efficient functions cannot be derived from other sampling algorithms. Further research in this field may yield such efficient functions from other sampling algorithms too.

Other reasons for choosing the Knuth-Yao algorithm are its efficiency and low entropy consumption. The Knuth-Yao and CDT (or its variants) are very popular choices for implementation of lattice-based cryptographic schemes due to their very efficient performance across different platforms. Howe et al. [35] has also recommended Knuth-Yao and CDT for constant-time discrete Gaussian sampling.

3.1 Our Observation: Mapping Random Bits to Samples

As explained in Section 2.2.2, the Knuth-Yao sampling is a random walk that starts from the root of a DDG tree and terminates when it hits a terminal node (Fig. 1). At each node, a random bit is used to select the sub-tree which will be explored in the next step. Hence, the path from the root of the tree to each terminal node is determined by a unique bit string. As each terminal node corresponds to a sample in the sample space, there exists a mapping from the set of random bit strings to the sample space.

Clearly, this mapping is many-to-one (as shown in Fig. 2). For example in Fig. 1, sample 0 is returned when the bit strings are 01 or 110. Or, if the random bits are extracted from random bit strings of length 4 then sample 0 is returned when the bit string is 01xx or 110x, where x can be either 0 or 1. Using the above observation, we can formulate the bits s_i of the sample s as a binary function of the random bit strings ($r = r_0 \cdots r_{n-1}$) as Eq. (1), assuming the samples can have maximum m bits and the probability matrix has n columns.

$$\begin{aligned} s_0 &= f^0(r_0, r_1, \dots, r_{n-1}) \\ s_1 &= f^1(r_0, r_1, \dots, r_{n-1}) \\ &\vdots \\ s_{m-1} &= f^{m-1}(r_0, r_1, \dots, r_{n-1}). \end{aligned} \quad (1)$$

To calculate a bit s_i of a sample, the respective binary expression f^i applies the corresponding set of binary operators on the input random bit string ($r_0 r_1 \cdots r_{n-1}$) irrespective of its value. Hence, for any i th bit s_i of the sample, the computation time t_i is always the same for any random input bit-string r . As an illustration, the toy distribution given in Fig. 1, the sample space has only 4 samples, hence $m = 2$ and $n = 4$. All possible input of bit-string of length $n = 4$

TABLE 1
Input-Output Table for the Arbitrary Distribution in Fig. 1

r_0	r_1	r_2	r_3	s_1	s_0	r_0	r_1	r_2	r_3	s_1	s_0
0	0	0	0	0	1	1	0	0	0	1	1
0	0	0	1	0	1	1	0	0	1	1	1
0	0	1	0	0	1	1	0	1	0	1	0
0	0	1	1	0	1	1	0	1	1	1	0
0	1	0	0	0	0	1	1	0	0	0	0
0	1	0	1	0	0	1	1	0	1	0	0
0	1	1	0	0	0	1	1	1	0	1	0
0	1	1	1	0	0	1	1	1	1	0	1

bits and their corresponding output sample for this distribution are shown in Table 1.

Using Karnaugh map minimization, we can calculate the sample bits (s_0, s_1) as,

$$\begin{aligned} s_0 &= \bar{r}_0 \bar{r}_1 \vee r_0 \bar{r}_1 \bar{r}_2 \vee r_0 r_1 r_2 r_3 \\ s_1 &= r_0 \bar{r}_1 \vee r_0 r_1 r_2 \bar{r}_3. \end{aligned}$$

It is worth noting that to ensure constant-time sampling, all the binary operators should be fully applied in the order specified in Eq. (1) for each sample regardless of the input random bit string. We use a bit-slicing methodology and bit-wise operators for this purpose. This will be explained in Section 3.3.

Each sample can thus be generated in *constant-time* by computing each of its bits in constant-time. Unlike other Gaussian sampling methods, this method neither requires a large precomputed table nor expensive computations. However, this method requires a *larger program memory* to store the formulae f^i .

3.2 Batching the Sampling Process

Bit-slicing is a Single Instruction Multiple Data (SIMD) operation to improve the efficiency of a programme by exploiting data level parallelism. Starting from Biham's implementation of DES [39], cryptographers have been using this method to speed up the execution of their algorithms for long time. Also, implementations using bit-slicing offers some immunity against side-channel attacks. Earlier, this method has been used for fast and side-channel secure AES implementations [40], [41].

In this section, we describe a method to speed up our sampling by generating multiple samples in a batch using bit-slicing. We utilize bit-wise Boolean operators on full processor words to achieve this. As shown in Eq. (1), each sample bit can be written as a function of n random bits. In the simplest approach, we can store the n random bits in $\lceil n/w \rceil$ variables, where w is the word length of the processor and compute the sample bits by extracting the random bits from variables as required. This way of generating sample bits is very inefficient and has a very low throughput.

However, using an efficient storage of random bits we can greatly improve the throughput. Let's assume, we want to generate k samples. So according to Eq. (1), we need nk bits. To store these bits efficiently such that we can use bit-slicing, we take n variables each of which stores k bits. The bit $j \in [0, k-1]$ of the i th variable $i \in [0, n-1]$ represents the input random bit r_i for sample j th as in Eq. (1). In other words, the i th variable stores all the i th input random bits

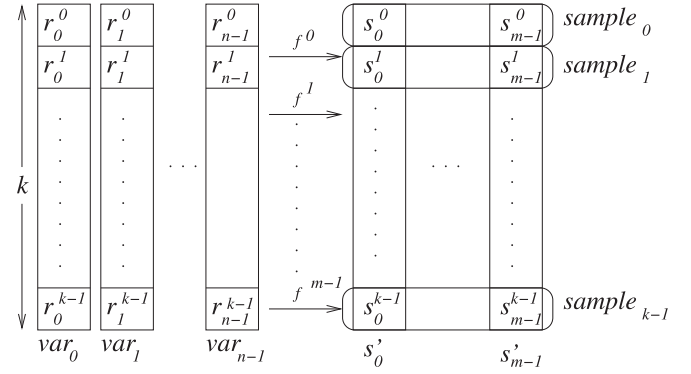


Fig. 3. Efficient storage of random bits and sample bits for bit-slicing. Here r_i^j represents i th input random bit of j th sample. Similarly, s_i^j represents i th output bit of j th sample.

to generate k samples (Fig. 3). We can then apply the bit-wise operators on these variables as indicated by Eq. (1). Alternatively, we can rewrite Eq. (1) as

$$\begin{aligned} s'_0 &= f^0(var_0, var_1, \dots, var_{n-1}) \\ s'_1 &= f^1(var_0, var_1, \dots, var_{n-1}) \\ &\vdots \\ s'_{m-1} &= f^{m-1}(var_0, var_1, \dots, var_{n-1}), \end{aligned} \quad (2)$$

Where each variable s'_t contains the t th sample bits s_t , $t \in [0, m-1]$ of the k samples. These variables are then used to extract the output sample bits to construct k samples. Here, evidently the maximum value of k is the word size w of the processor.

Therefore, using bit-wise Boolean operations and efficiently organizing the storage of input random bits, we can generate w samples simultaneously. This is explained in Fig. 3.

In Table 2, we show the performance of our bit-sliced sampler ($\sigma \approx 6.15543$) on an 2.66 GHz Intel i7-6600 Broadwell processor ($w = 64$) for different levels of precision

3.3 Analyzing the Constant-Time Execution

We ensure constant-time execution of the Boolean expressions by executing them till the end, i.e., not terminating them early depending on the data. Consider the following toy example where we compute

$$o = (a) \wedge (b \vee c). \quad (3)$$

Here a , b and c are random bits. If Eq. (3) is computed on a hypothetical 1-bit machine, then a smart compiler could minimize the computation time by ignoring the computation of $(b \vee c)$ whenever $a = 0$. In that case the implementation would not be constant-time.

TABLE 2
Time to Generate Samples for Different Precision

Precision λ	64	96	128
Probability of not finding a sample	2^{-59}	2^{-91}	2^{-122}
Clock-cycles			
(excluding random number generation)	4543	7548	11814

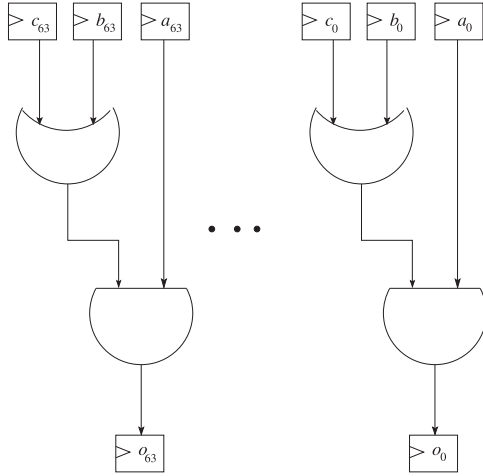


Fig. 4. Circuit-level abstraction of $O = (A) \wedge (B \vee C)$.

Now consider a bit-sliced implementation (as explained in Section 3.2) on a 64-bit processor. Here, 64-bit random integers A , B , and C can be considered as vectors of 64 random bits. Hence the computation of $O = (A) \wedge (B \vee C)$ using 64-bit ‘bitwise’ operations would compute all the 64 bits of O in parallel. Computation of $B \vee C$ becomes unnecessary only when all the bits of A become zero, which happens with probability 2^{-64} . A smart compiler would never insert data-dependent branches inside bitwise operations as it would be counterproductive. Nevertheless, a developer can even write the code in assembly to ensure no data-dependent branching by compiler.

Second, we analyze the timing behaviour of our approach at the circuit-level abstraction. In Fig. 4 we show a register-gate-level architecture diagram for computing $O = (A) \wedge (B \vee C)$. The output of any i th AND gate would settle early if its input a_i is zero. But the register o_i that stores the final result will get updated only after the next clock-edge transition. Since clock frequency of gate-level circuits is determined by the worst case propagation delay corresponding to the longest critical path, all of o_i are updated simultaneously and independently of input data. Hence, our approach is constant-time at the algorithm-level, instruction-level, and register-transfer-level, but it is non-constant-time at the gate-level and transistor-level.

4 PERFORMANCE AND COMPARISON

In this section, we compare our method with the CDT based constant-time algorithm using a C implementation. For the performance measurement, we use a discrete Gaussian distribution with standard deviation $\sigma \approx 6.15543$. In the next section, we justify our choice of this standard deviation.

4.1 Splitting the Gaussian Distribution

The BLISS-I [8] signature scheme uses a standard deviation $\sigma = 215$. However, as memory requirement to store the pre-computed table increases with increase in σ , sampling from a Gaussian distribution with such a large standard deviation is difficult due to large memory requirement. Also, due to the large precomputed tables, generating samples securely is highly inefficient. Pöppelman et al. [36] described a method to split this large standard deviation into two Gaussian

distributions with smaller standard deviation and later combining them to create a distribution with large standard deviation. They used Kullback-Liebler divergence, which is Rényi divergence of order 1 [34] instead of the more usual notion of statistical distance to show that the distribution created in this way is very close to the actual distribution. We extend their work by splitting the Gaussian distribution further in 4 smaller distributions.

We discuss the method by Pöppelman et al. very briefly here. To generate a sample $x \leftarrow D_\sigma$, two samples $x_1, x_2 \leftarrow D_{\sigma_1}$ are generated, and combined as $x_1 + k_1 x_2$. The σ, σ_1 and k_1 are related as $\sigma_1 = \frac{\sigma}{\sqrt{1+k_1^2}}$, for $\sigma = 215$, $k_1 = 11$ and $\sigma_1 \approx 19.5$. The Kullback-Leibler divergence of the sampled data created in this way from the actual distribution is $\leq 2^{-128}$. We split the standard deviation one more level. We split σ_1 such that $\sigma_2 = \frac{\sigma_1}{\sqrt{1+k_2^2}}$. Consequently, to generate a

sample $x \leftarrow D_\sigma$ we generate 4 samples $x_1, x_2, x_3, x_4 \leftarrow D_{\sigma_2}$ and combine them as $x = (x_1 + k_2 x_2) + k_1(x_3 + k_2 x_4)$.

We experimentally checked that setting $\tau_1 = 14$ and $k_2 = 3$ produces a Gaussian distribution with $\sigma_2 \approx 6.15543$ which has the desired divergence i.e., less than 2^{-128} from a Gaussian distribution with $\sigma = 215$.

For more details on closeness measures and their impact on security, we refer the reader to [24], [36], [42].

4.2 Performance

Our fully constant time sampling algorithm i.e., precision $\lambda = 128$, takes 11814 clock cycles to generate 64 samples from $\sigma \approx 6.15543$ on an Intel i7-6600 Broadwell processor running at 2.66 GHz. Hence, to generate 64 samples from the Gaussian distribution with $\sigma = 215$, used in BLISS-I [8], we need $4 \times 64 = 256$ samples from $\sigma \approx 6.15543$ i.e., $11814 \times 4 = 47256$ clock cycles or approximately 738 clock cycles to generate a single sample. If we include the cost to generate the pseudo-random numbers using SHAKE-128 standardized in FIPS-202 [43], our algorithm takes 26214 clock cycles to generate 64 samples with $\lambda = 128$ and $\sigma \approx 6.15543$ or approximately 1638 clock cycles to generate a single sample with $\sigma = 215$. Our high level implementation in C is only optimized by -O3 optimization of gcc. For efficiency, the Boolean functions f^i in Section 3.1 used to generate samples should be minimized. We used the logic minimization tool ESPRESSO [44] for this purpose.

As mentioned in Section 2.3, non-constant time methods can be converted to timing-attack resistant sampling methods using different countermeasures, which sacrifice their efficiency for security. Also, Howe et al. [35] has compared and analyzed such constant-time instantiations of different sampling algorithms. Their work shows that Knuth-Yao sampling with shuffle and constant-time cumulative distribution table (CDT) based methods are the most efficient for constant-time sampling. In this section, we compare our method with two of these methods for a similar probability of leakage.

The constant-time CDT sampler accesses all the elements of the CDF table for each sample. However, for a fair comparison with our method, we provide performance of CDT sampler for different levels of precision. For precision levels smaller than 128 instead of accessing the full table for each sample, we let the sampling method scan a part of the table

TABLE 3
Comparison of Clock Cycles for Different Constant Time Sampling with Similar Probability of Leakage for $\sigma \approx 6.15543$ to Generate 64 Samples on a 2.66 GHz Intel i7-6600 Broadwell Processor Using Only One Core

Algorithm		Time (in clockcycles)		
		$\lambda = 64$	$\lambda = 96$	$\lambda = 128$
Excluding pseudo-random number generation	CDT algorithm	11092	—	28231
	Full table KY scan [†]	11048	19683	31775
	Our algorithm	4543	7548	11814
Including pseudo-random number generation	CDT algorithm	19160	—	42181
	Full table KY scan [†]	11150	19836	31979
	Our algorithm	12660	17783	26214

SHAKE-128 extended output function has been used for pseudo-random number generation. [†] Time to generate one sample only.

corresponding to the level of precision. For instance, in our previous example Gaussian distribution with $\sigma \approx 6.15543$ precision $\lambda = 64$, we search only in an interval of $[0 - 39]$ or $[0 - 6.5\sigma]$. Since the CDT method performs comparisons between the random string and the table entries, we use either 64-bit or 128-bit comparisons taking into account the 64-bit word length of the processor.

As mentioned in Section 3.1, it is also possible to instantiate a constant-time Knuth-Yao sampling by accessing all the elements of a probability matrix for each sample using the bit-scanning algorithm described in [28]. But, the performance of this method is very poor as to generate a single sample it has to scan a large number of table entries. For example, in the case of our Gaussian distribution with $\sigma \approx 6.15543$, this method has to scan 6648 table entries to produce a single sample. We refer to this method as ‘Full Table Scan KY’.

We implemented all these methods in the C and compiled with -O3 flag in gcc-5.4 on a laptop with intel core-i7-Broadwell processor running at 2.66 GHz with turboboost and hyperthreading disabled and using only one core. The results are shown in Table 3.

The Knuth-Yao sampler with shuffling proposed by Roy et al. [31] is another method to prevent information leakage from the sampler. The method is described briefly in Section 2.3. To compare it with our method for a similar probability of leakage with our sampler with $\lambda = 64$, we need $k = 64$ which requires an enormous $O(2^{64})$ memory and a massive overhead for linear searching the table. Moreover, Bruinderink et al. [29] suggest that this method only increases the complexity of their attack. Peter Pessl [32] exploited this weakness of the sampler to break the BLISS signature scheme with an increased number of signatures.

4.2.1 Note on Memory Requirement

In our implementation the Boolean expressions are stored in the program memory. To implement the base sampler with 128-bit precision, the Boolean equations have total 18K AND, 5K OR and 11K NOT operations. Since this method does not require any table, no data memory is spent for storing tables. In our implementation we generate all the

random bits together and store them in 128 integers of size 64 bits. But it is possible to reduce the data memory to only 7 integers by generating the randomness on demand.

4.2.2 Results Using SIMD Instructions

New generation of Intel (starting with Sandy Bridge) and AMD (starting with Bulldozer) provides support for Advanced Vector Instructions (AVX). These instructions are an extension of the x86 instruction set architecture and facilitates SIMD processing on data of width up to 128 bits. Later, starting with Haswell processors, Intel introduced the AVX2 instruction set which increased this bit width to 256 bits. As described in Section 3.2, the throughput of our sampling algorithm can be readily increased by using AVX2 instruction sets. In Table 4, we report the time taken by our sampler to produce 256 samples at a time using AVX2 instructions. These results show almost 2x speed-up of our sampling algorithm using AVX2 instructions compared to the implementation on 64 bit processor.

4.2.3 Results in Hardware

To evaluate the performance of the proposed constant-time sampling algorithm in hardware, we designed two different architectures. The random bits are generated by an external source. We used SHAKE-128 for this purpose.

Bitsliced Knuth-Yao Sampler. A software-styled bitsliced architecture is shown in Fig. 5. The architecture is composed of an arithmetic unit, a register bank, a sample bank and a

TABLE 4
Time to Generate 256 Samples Using AVX2 Instructions

Precision λ	64	96	128
Clock-cycles (excluding random number generation)	6478	11565	19605




Fig. 5. Bitsliced hardware architecture for constant-time sampling.

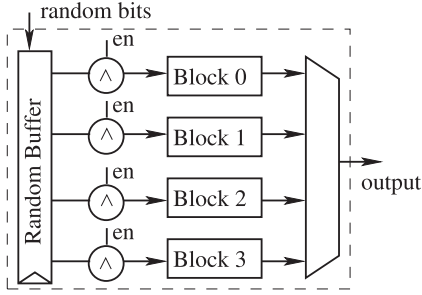


Fig. 6. Single-cycle architecture for constant-time sampling.

control unit. The arithmetic unit does bitwise AND, OR and NOT operations. The register bank stores the random numbers and the intermediate results. A naive implementation would require 112 registers to store all the random numbers for $\lambda = 112$. We observed that during evaluation of the Boolean expressions, the random numbers are used in a sequence and every random number is used only for a certain window of time. We used this observation to reduce the number of registers to only seven: a register is loaded with a random number as per requirement and once there is no more use of the random number, it is replaced by a new random number. The register bank contains additional five registers to store the intermediate results. The sample bank stores the computed sample values. The program memory is an addressable ROM where each entry contains an instruction code, operand-1 register id, operand-2 register id and the destination register id. The program memory was generated using a parser script from the Boolean expressions generated by ESPRESSO. When the width of the data path is set to 64 bits, the architecture of Fig. 5 requires 27,344 cycles (excluding the cost of random number generation) to compute 64 samples in a batch. We implemented the architecture on a Xilinx Virtex-6 FPGA *xc6vxc75t-2ff484*. As per place-and-route report, the architecture consumes 1,237 slice registers, 1,024 slice LUTs, and 15 RAMB36E1s (for the program ROM), and has a critical path delay of 7.5 ns. The number of cycles can be reduced by unrolling the program-instructions and keeping multiple arithmetic units in the architecture.

Decomposed Single Cycle Knuth-Yao Sampler. The architecture is shown in Fig. 6. The 112-bit register *random buffer* stores the input random bits. After *random buffer* is filled with the input random bits, the *enable* signal becomes true, and then the random bits are processed by the parallel combinational circuits *Block-0* to *Block-3*. These circuits implement the Boolean expressions of the sample-bits. In our architecture, *Block-0*, *Block1*, *Block2* and *Block-3* compute on the random bits of *random buffer* with the index ranges 0-31, 27-58, 54-86 and 80-111 respectively. All of these blocks compute in parallel and the output of the sampling operation is selected using the output-multiplexer. Hence excluding the cost of random number generation, only one cycle is spent by the parallel blocks to generate one sample.

We implemented the architecture (Fig. 6) on a Xilinx Virtex-6 FPGA *xc6vxc75t-2ff484*. As per place-and-route report, the architecture consumes 997 slice registers and 2,682 slice LUTs and has a critical path delay of 4.9ns.

Authorized licensed use limited to: BEIJING INSTITUTE OF TECHNOLOGY. Downloaded on April 12, 2024 at 08:45:37 UTC from IEEE Xplore. Restrictions apply.

5 EVALUATION

The implementation from Section 3 follows best-practice guidelines for constant-time code: constant program flow (no conditional branches), no secret-dependent memory accesses, and no usage of integer division nor multiplication operations. However, the fact that the high-level code looks constant time is no guarantee for the actual execution being constant time. Any piece in the tool chain may introduce a source of timing variability: in an extreme case, a very clever compiler would substitute the whole constant-time sampler with a faster, non constant-time one. Compilers and COTS architectures are currently designed to optimize for speed, code size, energy or power, but not security.

Thus, we resort to actual measurements to evaluate whether the resulting executable code runs in constant time on our platform or not. The evaluation of this section is empirical in nature and thus is bounded to the specific architecture, compiler and platform used.

5.1 Methodology

To assess timing variability we use leakage detection tests. Leakage detection tests were introduced by Coron, Naccache and Kocher [45], [46] shortly after the introduction of DPA [47] and were targeted towards hardware side-channel evaluations. Nowadays, this technique has been proven to be useful also for timing variability evaluation. In this section, we follow the methodology and test code from [48].

5.1.1 Leakage Detection for PRNGs

Generally speaking, we want to assess whether or not an adversary gets any advantage in distinguishing output samples from timing side-channel information. For that, we will deploy timing leakage detection tests to detect dependency between the execution time of the sampling procedure and input value to the Gaussian sampler. If the test fails to detect any dependency, the implementation is deemed secure. Note that the opposite outcome (there is detected leakage) is a necessary, but not sufficient, condition for an attack to work.

We design the timing leakage detection test as follows. We define two classes based on the input seed to the Gaussian sampler (that is, the input seed is treated as secret, and we aspire to detect any leakage dependent on this secret value). The two classes are defined as this: one class corresponds to a fix seed value; the other class is defined as a random seed value (fix-vs-random test). This choice, in contrast to a fix-vs-fix test, is expected to capture a broad set of leakages [49].

5.2 Platform

We perform the following experiments on the same platform from Section 4.2. We note that cycle counts are performed with the high-resolution Time Stamp Counter (RDTSC instruction).

The first implementation is the constant-time variant of Section 3.2. This version does not early abort and is meant to be constant time by design. We carry the evaluation to confirm that this is actually the case, i.e., the compiler or any other micro-architectural components do not introduce any source of timing variability.

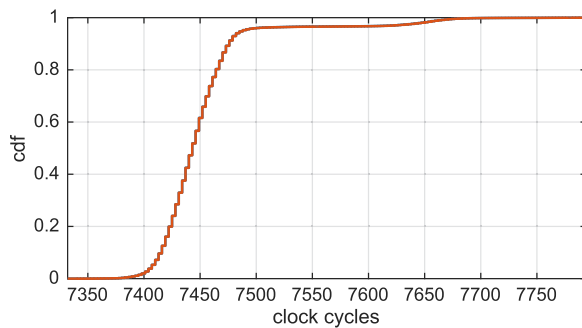


Fig. 7. Timing distribution cdfs for two classes in a fix-vs-random timing leakage detection test.

In Fig. 7 we plot the empirical cumulative distribution functions for both timing distributions, corresponding to the two classes on input values (fix or random). The two distributions are actually indistinguishable and their cdfs overlap. We can see that the distributions are centered around 7450 cycles and there is a small class-independent variability (≈ 100 cycles). This measurement noise could be caused by spurious interruptions by the operating system, or by the processor itself (for example, due to branch mispredictions). The leakage detection t-test statistic does not surpass the threshold of ± 4.5 and hence does not detect any leakage with up to 6 million iterations of the bitsliced sampling process. Various pre-processing options were explored with identical results.

We also perform a Kolmogorov-Smirnov (KS) test. The advantage is that it may detect that two distributions are different even if they share the same mean. The value of the statistic is 0.000625 which is lower than the cutoff value 0.001282. Thus, the KS test cannot reject the null hypothesis that both distributions are identical.

6 CONCLUSION AND DISCUSSION

In this paper, we present a constant-time version of the Knuth-Yao sampling algorithm. We also perform optimization it to make the sampling algorithm many times faster than existing leakage resistant discrete Gaussian sampling algorithms. These optimizations do not require any special hardware and can be implemented on most modern processors.

We are aware that though this method does not require *large data memory* to store large precomputed tables, it requires a *larger program memory* than other methods, which is not so much problem for desktop computers as it is for devices with very limited resources. Future research will try to reduce the program memory by possibly tweaking the minimization procedure of Boolean functions or devising encoding schemes to reduce the storage of program memory. These methods may sacrifice its efficiency to some extent but will be suitable for devices with limited resources.

There are some simple optimizations that could be applied to make the method more efficient. For example, as our method does not require frequent external-memory accesses and has a high degree of parallelism, it can be exploited to design fast constant-time discrete Gaussian sampling on multi-core processors. Also, minimizations of the Boolean functions has a direct impact on the efficiency

of the sampling algorithm. Our current minimization using ESPRESSO is not optimal and uses only bitwise AND, OR and NOT operations. There is a possibility to use different tools to get a better minimization of the Boolean functions, which will immediately translate into a faster sampling process. Also, we can see from our results in Table 3 that generating pseudo-random numbers using SHAKE-128 takes a significant portion of the running time of the whole sampling operation. It will be interesting to test the performance of the sampler using different pseudo-random number generators. We leave this for further research.

In this work our focus was a proof of concept implementation and we have not focused on the optimization of the Boolean equations. In the future we would investigate efficient encoding techniques for optimizing the memory requirement.

Comparison with [42]. A concurrent work by Micciancio and Walter [42] proposes an algorithm that can sample from a discrete Gaussian distribution with arbitrary standard deviation. The algorithm uses convolution in a recursive fashion to generate samples from the target distribution by combining several samples from a fixed and small ‘base’ distribution. The technique is generic because it can sample from discrete Gaussian distributions with arbitrary standard deviations and centers. If the base sampler is constant time then sampling from the target distribution can be performed in constant time. However, the authors make no effort to make the base sampling operation constant-time and remarks that substantial amount of design and implementation effort would be required to make their algorithm constant-time (see Section 6.3 of [42]). In our work we show a methodology to perform sampling from the ‘base’ distribution in constant time. In our opinion these two works complement each other and [42] could use our methodology to achieve truly constant-time discrete Gaussian sampling.

ACKNOWLEDGMENTS

This work was supported in part by the Research Council KU Leuven: C16/15/058. In addition, this work was supported in part by the Flemish Government, by the Hercules Foundation AKUL/11/19, and by the European Commission through ICT programme under contract through the Horizon 2020 research and innovation programme under contract No H2020-ICT-2014-644371 WITDOM, H2020-ICT-2014-644209 HEAT, H2020-ICT-2014-645622 PQCRYPTO and Cathedral ERC Advanced Grant 695305. Angshuman Karmakar was funded by Erasmus Mundus scholarship.

REFERENCES

- [1] W. Diffie and M. Hellman, “New directions in cryptography,” *IEEE Trans. Inf. Theory*, vol. 22, no. 6, pp. 644–654, Nov. 1976.
- [2] P. W. Shor, “Polynomial time algorithms for prime factorization and discrete logarithms on a quantum computer,” *SIAM J. Sci. Comput.*, vol. 26, 1997, Art. no. 1484.
- [3] J. Proos and C. Zalka, “Shor’s discrete logarithm quantum algorithm for elliptic curves,” arXiv:quant-ph/0301141, Jan. 2003.
- [4] O. Regev, “New lattice-based cryptographic constructions,” *J. ACM (JACM)*, New York, NY, USA: ACM, vol. 51, no. 6, pp. 899–942, Nov. 2004. [Online]. Available: <https://dl.acm.org/citation.cfm?id=1039490>
- [5] M. Ajtai, “Generating hard instances of lattice problems (extended abstract),” in *Proc. 28th Annu. ACM Symp. Theory Comput.*, 1996, pp. 99–108. [Online]. Available: <http://doi.acm.org/10.1145/237814.237838>

- [6] V. Lyubashevsky, C. Peikert, and O. Regev, "On ideal lattices and learning with errors over rings," in *Proc. Advances Cryptology 29th Annu. Int. Conf. Theory Appl. Cryptographic Techn.*, 2010, pp. 1–23. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-13190-5_1
- [7] D. Micciancio, "Generalized compact knapsacks, cyclic lattices, and efficient one-way functions," *Comput. Complexity*, vol. 16, no. 4, pp. 365–411, 2007. [Online]. Available: <http://dx.doi.org/10.1007/s00037-007-0234-9>
- [8] L. Ducas, A. Durmus, T. Lepoint, and V. Lyubashevsky, "Lattice signatures and bimodal gaussians," in *Proc. Advances Cryptology 33rd Annu. Cryptology Conf.*, 2013, pp. 40–56. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-40041-4_3
- [9] P. S. L. M. Barreto, P. Longa, M. Naehrig, J. E. Ricardini, and G. Zanon, "Sharper ring-LWE signatures," *Cryptology ePrint Archive*, Report 2016/1026, 2016. [Online]. Available: <http://eprint.iacr.org/2016/1026>
- [10] S. Akleyek, N. Bindel, J. Buchmann, J. Krämer, and G. A. Marson, "An efficient lattice-based signature scheme with provably secure instantiation," in *Proc. 8th Int. Conf. Cryptology in Africa*, 2016, pp. 44–60. [Online]. Available: http://dx.doi.org/10.1007/978-3-319-31517-1_3
- [11] R. Lindner and C. Peikert, "Better key sizes (and attacks) for lwe-based encryption," in *Proc. Cryptographers' Track at RSA Conf.*, 2011, pp. 319–339. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-19074-2_21
- [12] E. Alkim, L. Ducas, T. Pöppelmann, and P. Schwabe, "Post-quantum key exchange - a new hope," in *Proc. 25th USENIX Security Symp.*, 2016, pp. 327–343. [Online]. Available: https://www.usenix.org/system/files/conference/usenixsecurity16/sec16_paper_alkim.pdf
- [13] J. Bos, C. Costello, L. Ducas, I. Mironov, M. Naehrig, V. Nikolaenko, A. Raghunathan, and D. Stebila, "Frodo: Take off the ring! practical, quantum-secure key exchange from LWE," in *Proc. ACM SIGSAC Conf. Comput. Commun. Security*, 2016, pp. 1006–1018. [Online]. Available: <http://doi.acm.org/10.1145/2976749.2978425>
- [14] C. Gentry, C. Peikert, and V. Vaikuntanathan, "Trapdoors for hard lattices and new cryptographic constructions," in *Proc. 40th Annu. ACM Symp. Theory Comput.*, 2008, pp. 197–206. [Online]. Available: <http://doi.acm.org/10.1145/1374376.1374407>
- [15] D. Cash, D. Hofheinz, E. Kiltz, and C. Peikert, "Bonsai trees, or how to delegate a lattice basis," in *Proc. 29th Annu. Int. Conf. Theory Appl. Cryptographic Techn.*, 2010, pp. 523–552. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-13190-5_27
- [16] S. Agrawal, D. Boneh, and K. Boyen, "Efficient lattice (h)ibe in the standard model," in *Proc. 29th Annu. Int. Conf. Theory Appl. Cryptographic Techn.*, 2010, pp. 553–572. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-13190-5_28
- [17] S. "Agrawal, D. Boneh, and X. Boyen, "Lattice basis delegation in fixed dimension and shorter-ciphertext hierarchical IBE," in *Proc. 30th Annu. Conf. Advances Cryptology*, 2010, pp. 98–115. [Online]. Available: <http://dl.acm.org/citation.cfm?id=1881412.1881420>
- [18] Z. Brakerski and V. Vaikuntanathan, "Efficient fully homomorphic encryption from (standard) LWE," in *Proc. IEEE 52nd Annu. Symp. Foundations Comput. Sci.*, Oct. 2011, pp. 97–106.
- [19] Z. Brakerski, C. Gentry, and V. Vaikuntanathan, "(leveled) fully homomorphic encryption without bootstrapping," *ACM Trans. Comput. Theory*, vol. 6, no. 3, pp. 13:1–13:36, Jul. 2014. [Online]. Available: <http://doi.acm.org/10.1145/2633600>
- [20] Z. Brakerski, "Fully homomorphic encryption without modulus switching from classical gapsvp," in *Proc. 32nd Annu. Cryptology Conf.*, 2012, pp. 868–886. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-32009-5_50
- [21] C. Gentry, "A fully homomorphic encryption scheme," Ph.D. dissertation, Stanford Univ., Stanford, CA, USA, 2009, <https://crypto.stanford.edu/craig/>
- [22] V. Lyubashevsky, "Lattice signatures without trapdoors," in *Proc. 31st Annu. Int. Conf. Theory Appl. Cryptographic Techn.*, 2012, pp. 738–755. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-29011-4_43
- [23] P. Weiden, A. Hülsing, D. Cabarcas, and J. Buchmann, "Instantiating treeless signature schemes," *Cryptology ePrint Archive*, Report 2013/065, 2013, [Online]. Available: <https://eprint.iacr.org/2013/065.pdf>
- [24] C. Peikert, "An efficient and parallel gaussian sampler for lattices," in *Proc. 30th Annu. Cryptology Conf.*, 2010, pp. 80–97. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-14623-7_5
- [25] L. Ducas and P. Q. Nguyen, "Faster gaussian lattice sampling using lazy floating-point arithmetic," in *Proc. 18th Int. Conf. Theory Appl. Cryptology Inf. Security*, 2012, pp. 415–432. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-34961-4_26
- [26] J. Buchmann, D. Cabarcas, F. Göpfert, A. Hülsing, and P. Weiden, "Discrete ziggurat: A time-memory trade-off for sampling from a gaussian distribution over the integers," in *Revised Selected Papers on Selected Areas in Cryptography*. New York, NY, USA: Springer-Verlag, 2014, pp. 402–417. [Online]. Available: http://dx.doi.org/10.1007/978-3-662-43414-7_20
- [27] N. C. Dwarakanath and S. D. Galbraith, "Sampling from discrete gaussians for lattice-based cryptography on a constrained device," *Applicable Algebra Eng. Commun. Comput.*, vol. 25, no. 3, pp. 159–180, 2014. [Online]. Available: <http://dx.doi.org/10.1007/s00200-014-0218-3>
- [28] S. S. Roy, F. Vercauteren, and I. Verbauwhede, "High precision discrete gaussian sampling on fpgas," in *Revised Selected Papers on Selected Areas in Cryptography – SAC 2013 – Volume 8282*. New York, NY, USA: Springer, 2014, pp. 383–401. [Online]. Available: http://dx.doi.org/10.1007/978-3-662-43414-7_19
- [29] L. Groot Bruinderink, A. Hülsing, T. Lange, and Y. Yarom, "Flush, gauss, and reload – a cache attack on the bliss lattice-based signature scheme," in *Proc. 18th Int. Conf. Cryptographic Hardware Embedded Syst.*, 2016, pp. 323–345. [Online]. Available: http://dx.doi.org/10.1007/978-3-662-53140-2_16
- [30] A. K. Lenstra, H. W. Lenstra, and L. Lovász, "Factoring polynomials with rational coefficients," *Mathematische Annalen*, vol. 261, no. 4, pp. 515–534, Dec. 1982.
- [31] S. S. Roy, O. Reparaz, F. Vercauteren, and I. Verbauwhede, "Compact and side channel resistant discrete gaussian sampling," *Cryptology ePrint Archive*, Report 2014/591, 2014. [Online]. Available: <https://eprint.iacr.org/2014/591.pdf>
- [32] P. Pessl, "Analyzing the shuffling side-channel countermeasure for lattice-based signatures," in *Proc. 17th Int. Conf. Cryptology India*, 2016, pp. 153–170. [Online]. Available: http://dx.doi.org/10.1007/978-3-319-49890-4_9
- [33] J. W. Bos, C. Costello, M. Naehrig, and D. Stebila, "Post-quantum key exchange for the TLS protocol from the ring learning with errors problem," in *Proc. IEEE Symp. Security Privacy*, May 2015, pp. 553–570.
- [34] S. Bai, A. Langlois, T. Lepoint, D. Stehlé, and R. Steinfeld, "Improved security proofs in lattice-based cryptography: Using the Rényi divergence rather than the statistical distance," in *Proc. 21st Int. Conf. Theory Appl. Cryptology Inf. Security*, 2015, pp. 3–24. [Online]. Available: http://dx.doi.org/10.1007/978-3-662-48797-6_1
- [35] J. Howe, A. Khalid, C. Rafferty, F. Regazzoni, and M. O'Neill, "On practical discrete gaussian samplers for lattice-based cryptography," *IEEE Trans. Comput.*, vol. 67, no. 3, pp. 322–334, Mar. 2016, <https://ieeexplore.ieee.org/document/7792671/>
- [36] T. Pöppelmann, L. Ducas, and T. Güneysu, "Enhanced lattice-based signatures on reconfigurable hardware," in *Proc. 16th Int. Workshop Cryptographic Hardware Embedded Syst.*, 2014, pp. 353–370. [Online]. Available: http://dx.doi.org/10.1007/978-3-662-44709-3_20
- [37] C. Du and G. Bai, "Towards efficient discrete gaussian sampling for lattice-based cryptography," in *Proc. 25th Int. Conf. Field Programmable Logic Appl.*, Sep. 2015, pp. 1–6.
- [38] D. Knuth and A. Yao, "The complexity of nonuniform random number generation," in *Algorithms and Complexity: New Directions and Recent Results*. Cambridge, MA, USA: Academic Press, 1976.
- [39] E. Biham, "A fast new des implementation in software," in *Proc. 4th Int. Workshop Fast Softw. Encryption*, 1997, pp. 260–272. [Online]. Available: <http://dx.doi.org/10.1007/BFb0052352>
- [40] E. Käsper and P. Schwabe, "Faster and timing-attack resistant aes-gcm," in *Proc. 11th Int. Workshop Lausanne Cryptographic Hardware Embedded Syst.*, 2009, pp. 1–17. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-04138-9_1
- [41] C. Rebeiro, D. Selvakumar, and A. S. L. Devi, "Bitslice implementation of AES," in *Proc. 5th Int. Conf. Cryptology Netw. Security*, 2006, pp. 203–212. [Online]. Available: http://dx.doi.org/10.1007/11935070_14
- [42] D. Micciancio and M. Walter, "Gaussian sampling over the integers: Efficient, generic, constant-time," in *Proc. 37th Annu. Int. Cryptology Conf.*, 2017, pp. 455–485.
- [43] N. I. of Standards and T. 2015, "SHA-3 standard: Permutation-based hash and extendable-output functions," *Federal Inf. Process. Stds. PUB 202*, 2015. [Online]. Available: <http://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.202.pdf>

- [44] R. K. Brayton, A. L. Sangiovanni-Vincentelli, C. T. McMullen, and G. D. Hachtel, *Logic Minimization Algorithms VLSI Synthesis*. Norwell, MA, USA: Kluwer Academic Publishers, 1984.
- [45] J. Coron, P. C. Kocher, and D. Naccache, "Statistics and secret leakage," in *Proc. Int. Conf. Financial Cryptography*, 2000, pp. 157–173. [Online]. Available: http://dx.doi.org/10.1007/3-540-45472-1_12
- [46] J. Coron, D. Naccache, and P. C. Kocher, "Statistics and secret leakage," *ACM Trans. Embedded Comput. Syst.*, vol. 3, no. 3, pp. 492–508, 2004. [Online]. Available: <http://doi.acm.org/10.1145/1015047.1015050>
- [47] P. C. Kocher, J. Jaffe, and B. Jun, "Differential power analysis," in *Proc. 19th Annu. Int. Cryptology Conf.*, 1999, pp. 388–397. [Online]. Available: http://dx.doi.org/10.1007/3-540-48405-1_25
- [48] O. Reparaz, J. Balasch, and I. Verbauwhede, "Dude, is my code constant time?" in *Des. Autom. Test Eur. Conf. Exhib.*, 2017, Art. no. 14.
- [49] F. Durvaux and F. Standaert, "From improved leakage detection to the detection of points of interests in leakage traces," in *Proc. 35th Annu. Int. Conf. Theory Appl. Cryptographic Techn.*, 2016, pp. 240–262. [Online]. Available: https://doi.org/10.1007/978-3-662-49890-3_10



Angshuman Karmakar received the BE degree in computer science and engineering from Jadavpur University, Kolkata, and the MTech degree in computer science and engineering from the Indian Institute of Technology, Kharagpur. He is working toward the PhD degree in the Department of Electrical Engineering (ESAT), KU Leuven, Belgium. He has worked as an engineer in Citrix R&D India Ltd, Bangalore for two years. His research interests include broadly post-quantum cryptography, public-key cryptography and cryptanalysis.



Sujoy Sinha Roy received the MS degree in computer science and engineering from the Indian Institute of Technology Kharagpur, and the PhD degree in electrical engineering from the Katholieke Universiteit Leuven, Belgium. He is currently a post-doctoral researcher with the COSIC, Department of Electrical Engineering, Katholieke Universiteit Leuven, Belgium. His research area has been broadly in the field of efficient implementation of public key cryptography.



Oscar Reparez received the MSc degree in electrical engineering from ETSI Telecomunicacion, Madrid, Spain, and the PhD degree from Katholieke Universiteit Leuven. He is currently a post-doctoral researcher with the COSIC, Department of Electrical Engineering, Katholieke Universiteit Leuven, Belgium. His research focuses on applied cryptography.



Frederik Vercauteren received the MSc degree in computer science, the MSc degree in pure mathematics, and the PhD degree in electrical engineering from the Katholieke Universiteit Leuven, Belgium. He is a professor in the research group COSIC, Electrical Engineering Department, KU Leuven in Belgium. Previously, he was a lecturer with the Department of Computer Science, University of Bristol, United Kingdom. His research interests include applications of computational number theory and arithmetic geometry in cryptography, in particular post-quantum cryptography and homomorphic encryption.



Ingrid Verbauwhede is a professor in the research group COSIC, Electrical Engineering Department, KU Leuven, in Belgium. At COSIC, she leads the embedded systems and hardware group. She is also adjunct professor with the EE Department, University of California, Los Angeles, California. She is a member of the Royal Academy of Belgium for Science and the Arts. She is a recipient of an ERC Advanced Grant in 2016. She will receive the IEEE 2017 Computer Society Technical Achievement Award. She is a pioneer in the field of efficient and secure implementations of cryptographic algorithms in embedded context on ASIC, FPGA and embedded SW. She is the author and co-author of more than 300 publications at conferences, journals, book chapters and books. Her list of publications and patents is available at www.esat.kuleuven.be/cosic.

► For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.